

ICS 1019 – Knowledge Representation & Reasoning

30% Project

Joel Azzopardi

March 2020

This project will require you to implement 2 reasoning methods that were covered during lectures. These are:

1. The parsing of Horn Clauses in CNF, and reasoning with them using the back-chaining SLD procedure (covered in Chapter 5 of the Brachman and Levesque book and slides).
2. Construction of Inheritance Networks, and reasoning with them using pre-emption and redundancy (covered in Chapter 10 of the Brachman and Levesque book and slides).

1 Parsing of Horn Clauses, and Reasoning using Back-Chaining

This section is based on Chapter 5 of the Brachman and Levesque book. It involves the construction of a Knowledge Base from a list of Horn Clauses provided in Conjunctive Normal Form (CNF), and then performing reasoning on them using the Back-chaining algorithm.

1.1 Aim

The aim is to use any third generation language of your choice (Java, JavaScript, C, C++, C#, Python, ...) to implement a program that can read and parse clauses written in CNF and perform reasoning on the resulting knowledge base. The user can then query your program to check if a query clause can be resolved to an empty clause. These are described in more detail in the subsections below.

You are **not expected** to produce any fancy interfaces. A basic interface (even a command line interface) will suffice as long as the system allows the user inputs and provides the relevant outputs as described below. No marks are assigned to the design of the interface.

1.2 Input clauses

Your system needs to read the initial Knowledge Base from a file (in case of a command-line or GUI interface), or by reading it in its entirety from a textbox (only in case of a GUI interface - e.g. in JavaScript). Note that your system should **not** read the statements one by one from the command prompt.

The Knowledge Base needs to consist (**strictly**) of statements in the following format:

`[<literal 1>,<literal 2>, ...]`

Literals are separated by a ‘,’ and each clause starts on a new line. The input will have no extra spaces.

An example input can be:

```
[ Toddler ]
[ ¬Toddler , Child ]
[ ¬Child , ¬Male , Boy ]
[ ¬Infant , Child ]
[ ¬Child , ¬Female , Girl ]
[ Female ]
[ Male ]
```

If you encounter issues handling the \neg operator, you can use this symbol instead: ‘!’ as shown below:

```
[ Toddler ]  
[!Toddler , Child ]  
[!Child ,!Male , Boy ]  
[!Infant , Child ]  
[!Child ,!Female , Girl ]  
[ Female ]  
[ Male ]
```

You are to assume that the input clauses are all horn clauses and that they do not contain any errors. Your system does not need to try to validate the input (i.e. detect incorrect input), or be robust enough to cater for errors in the input.

Note also that the input clauses can **only** contain atomic literals (or their negation). You should not worry about handling functions or predicates.

1.3 User Queries

The user should be allowed to query the system by submitting his/her negated query (i.e. $\neg\alpha$) in CNF. Your system does not need to attempt to do the query negation automatically.

An example query to the above Knowledge Base is:

```
[  $\neg$ Girl ]
```

In case of issues handling the \neg operator, you can accept the following format instead:

```
[!Girl ]
```

1.4 Resolution

Your system should load the knowledge base and the user query (all in CNF). It should then use the back-chaining SLD procedure to try to resolve to the empty

clause ($[]$). This procedure is found in Page 92 of the book, and in Slide 87 of the corresponding slide deck. As output, the system only needs to output messages such as “**SOLVED**” if the empty clause was obtained through resolution and “**NOT SOLVED**” if not.

You can have your system output informative messages as well before the final ‘**SOLVED**’ or ‘**NOT SOLVED**’.

Note that the back-chaining SLD procedure uses recursion. You need to be careful from infinite loops when implementing it.

1.5 Execution

As mentioned before, your system should be executed in only one of the two modes listed below:

1. Command line – in this case the file should be passed as a command-line paramter as shown below:
 - `./backchaining.exe kb1.txt`
 - `python backchaining.py kb1.txt`
 - `java -jar backchaining.jar kb1.txt`
2. GUI – e.g. using Javascript, and the knowledge base should be pasted in its entirety.

Instructions on how to execute the program should be provided together with a sample input (knowledge-base and query that are known to work).

2 Construction and Reasoning with Inheritance Networks

This section is based on Chapter 10 of the Brachman and Levesque book. It involves the construction of simple inheritance networks, and then performing reasoning on these inheritance networks to resolve conflicts arising from multiple inheritance.

2.1 Aim

The aim is to use any third generation language of your choice (Java, JavaScript, C, C++, C#, Python, ...) to implement a program that can construct an inheritance network based on the input provided, and perform reasoning on the constructed network. The user can then query your program on possible inheritances. These are described in more detail in the subsections below.

You are not expected to produce any fancy interfaces. A basic interface (even a command line interface) will suffice as long as the system allows the user inputs and provides the relevant outputs as described below.

2.2 Input to construct Inheritance Network

Your system needs to read the initial Knowledge Base from a file (in case of a command-line or GUI interface), or by reading it in its entirety from a textbox (only in case of a GUI interface - e.g. in JavaScript). Note that your system should **not** read the statements one by one from the command prompt.

The format is shown below. Concept names will only have alpha-numeric characters in them (no spaces and/or special characters), and separate from the **IS-A** or **IS-NOT-A** operator by a single space. Each statement starts on a new line. The input will have no extra spaces or characters

```
<Sub-Concept> IS-A <Super-Concept>  
<Sub-Concept> IS-NOT-A <Super-Concept>
```




Figure 1: Example inheritance network

Given the example below, your program should construct the inheritance network shown in Figure 1.

```

Clyde IS-A FatRoyalElephant
FatRoyalElephant IS-A RoyalElephant
Clyde IS-A Elephant
RoyalElephant IS-A Elephant
RoyalElephant IS-NOT-A Gray
Elephant IS-A Gray

```

Note that you are **not** required to provide a graphical representation of the inheritance network.

You are to assume that concept names do not contain any spaces, and that the input provided does not contain any errors (your system does not need to try to understand incorrect input). You are also to assume that the inheritance network is **acyclic**. This means that your system does **not** need to perform any validation.

2.3 User Queries

The user should be allowed to query the system using only one type of command:

<Sub-Concept> **IS-A** <Super-Concept>

For example,

Clyde **IS-A** Gray

2.4 Query Handling

2.4.1 Searching for all possible paths

Upon receiving a query, the system should output all the possible paths from the query's sub-concept to the query's super-concept.

For the example above, the system should output the following paths:

- Clyde **IS-A** Elephant **IS-A** Gray
- Clyde **IS-A** FatRoyalElephant **IS-A** RoyalElephant **IS-A** Elephant **IS-A** Gray
- Clyde **IS-A** FatRoyalElephant **IS-A** RoyalElephant **IS-NOT-A** Gray

Remember that a path can only have a single **IS-NOT-A** link, and this can only occur at the very end. If whilst searching for the possible paths, your system encounters a **IS-NOT-A** link, it does not need to search further from the node at the end of this **IS-NOT-A** link.

2.4.2 Outputting the preferred path/s

The system should then provide which path is preferred according to the **shortest distance** metric, and according to the **inferential distance** metric. Note the

preferred path calculated using *inferential distance* should not include any **inadmissible links** – i.e. neither **redundant links** nor **pre-empted** links. More information about this may be found in slides 173 – 174

For the example above, these outputs will be:

- Preferred Path (Shortest Distance):
 - Clyde **IS-A** Elephant **IS-A** Gray
- Preferred Path (Inferential Distance):
 - Clyde **IS-A** FatRoyalElephant **IS-A** RoyalElephant **IS-NOT-A** Gray

If there is more than 1 preferred path (or the system can not determine the best path), all ‘preferred’ paths should be output.

2.5 Execution

As in the previous problem, your system should be executed in only one of the two modes listed below:

1. Command line – in this case the file should be passed as a command-line paramter as shown below:
 - `./inheritancenets.exe kb2.txt`
 - `python inheritancenets.py kb2.txt`
 - `java -jar inheritancenets.jar kb2.txt`
2. GUI – e.g. using Javascript, and the knowledge base should be pasted in its entirety.

Instructions on how to execute the program should be provided together with a sample input (knowledge-base and query that are known to work).

3 Documentation / Report

You are to write a short documentation (about 2-4 pages per problem) that describes your system, the data structures used, and how the operations were handled. Your report should not exceed 9 pages in total.

Most importantly, your documentation should include the following (for each of the two problems):

- Specific (short) instructions on how to execute your program/s. The instructions should **not** direct the user to use NetBeans, Eclipse or any other IDE. If you use Java, you should provide instructions that are based on the *java* command, and you should provide jar file/s. You should also provide the version of the programming language you used for implementation. (e.g. Python 2.7, Python 3, ...).
- Data examples which are known to work – the input knowledge base as well as the query. The input knowledge base should be provided as a text file with the submission.
- Specification of which parts of the project have been implemented and are working successfully; the parts that are only partially working; those are not working; and those that have not been attempted.

Marks will be given accordingly, but effort for parts that are not or are partially working will also be taken into consideration.

Note that failure to provide such instructions will result in loss of marks attributed to them, as well as potentially having your program considered as not working (if it does not work with the input cases provided upon testing it for marking).

4 Mark Distribution

- Back-Chaining on Horn Clauses

- Construction of knowledge base according to user inputs – 15%
 - Resolution by backchaining – 20%.
 - Documentation – 7%.
 - Execution instructions + input data examples + specification of completion – 8%.
- Inheritance Networks
 - Construction of inheritance network according to user inputs – 6%
 - Output of possible paths according to user queries – 12%.
 - Determination of shortest path – 5%.
 - Determination of preferred path according to inferential distance – 12%.
 - Documentation – 7%.
 - Execution instructions + input data examples + specification of completion – 8%.

5 Group Work

Note that this project is an individual project. No group work is permitted and plagiarism will be penalised.

6 Deadline and Submission

The deadline for this project is **Monday 11th May, 2020** at **noon**. The project should be uploaded on VLE. There will be 2 upload areas, one for the documentation (to be uploaded as *doc*, *docx* or *pdf*), and one for the program sources and example data (to be uploaded as *zip*, *rar*, *tar.gz* or *tgz*).

7 Difficulties

In case of difficulties, please do not hesitate to send an email on *joel.azzopardi@um.edu.mt*.

It is recommended that you are familiar with the relevant concepts from the lecture notes, and that you actually use the algorithms provided in the book/slides.